

Diabetes Onset Prediction and Risk Stratification Using an Optimized Multilayer Perceptron (MLP) and Conventional Machine Learning Baselines

Course Faculty: Dr. Gyana Ranjan Patra | Institution: Centre for AI & ML / Dept. of CSE | Term: 2026

Project Investigators & Contributors:

Sl. No.	Student Name	Registration Number	Department / Role
1	Tribhuwan Singh	2341019538	CSE / ML Lead & Pipeline Architect
2	Surajit Sahoo	2341019165	CSE / Preprocessing & Imputation
3	Anwesha Srichandan	2341019594	CSE / Feature Engineering & EDA
4	Priti Rani Maity	2341013065	CSE / Evaluation & Benchmark Analytics

1. Executive Summary

Diabetes mellitus represents one of the most pressing global metabolic health crises. Early diagnosis is vital to prevent severe secondary microvascular and macrovascular complications. This research project develops an end-to-end, data-leakage-free machine learning system to predict diabetes onset from physiological and demographic markers using the Pima Indians Diabetes Database (N = 768). The workflow resolves biologically implausible zero values in physiological attributes (Glucose, Blood Pressure, Skin Thickness, Insulin, BMI) via training-fitted median imputation, expands raw features into a 24-dimensional domain representation, and benchmarks seven conventional classification baselines against an optimized Scikit-Learn Multilayer Perceptron (MLPClassifier). The final tuned MLP architecture (24 -> 64 -> 32 -> 1, ReLU activation, Adam optimizer, alpha=0.001, batch size=32) achieved an untouched test set Accuracy of 82.76%, Sensitivity of 75.00%, Specificity of 86.84%, F1-Score of 75.00%, and ROC-AUC of 0.8681, successfully demonstrating neural pattern extraction on tabular biometric data.

2. Motivation & Clinical Importance

2.1 Background: Over 537 million adults globally live with diabetes, a figure projected to reach 783 million by 2045. Early screening is crucial because clinical onset often remains asymptomatic until irreversible metabolic damage occurs.

2.2 Limitations of Existing Tools: Traditional clinical risk calculators use static univariate cutoffs that fail to capture complex non-linear interactions between adiposity, insulin resistance, and pancreatic beta-cell compensation.

2.3 Need for Research: Many existing ML studies suffer from severe data leakage (fitting imputers/scalers globally prior to splitting) and mishandle physiological zero readings as true numerical values. This research establishes a leak-free, reproducible neural pipeline.

3. Problem Statement & Research Gap

Current Situation	High global diabetes prevalence; static univariate cutoffs miss complex metabolic interactions.
Research Gap	Prior ML studies often exhibit data leakage, ignore zero-value biological artifacts, and lack multi-algorithm benchmarks against neural MLPs.
Proposed Solution	An authoritative leak-free pipeline with median zero imputation, 24-dim domain engineering, systematic MLP grid search, 8-model benchmark, and cloud deployment.

4. Research Objectives (SMART) & Chapter Mapping

The research objectives follow the SMART criteria using precise action verbs:

Objective	Action Verb	Description	Chapter
O1	Analyze	Conduct rigorous EDA on feature distributions, skewness, and class imbalance (65.1% vs 34.9%).	Section 6
O2	Implement	Build leak-free zero-detection and training-fitted median imputation pipeline.	Section 7
O3	Design	Formulate 16 domain features (WHO BMI bins, ADA stages, metabolic ratios) -> 24 features.	Section 8
O4	Develop	Implement 6 baseline classifiers (LogReg, KNN, SVM, DT, RF, GBDT) & baseline MLP.	Section 9
O5	Optimize	Execute systematic Grid Search over MLP architectures, L2 penalties, and learning rates.	Section 10
O6	Evaluate	Benchmark all 8 models on untouched test partition (N=116) across 6 evaluation metrics.	Section 11
O7	Deploy	Construct and deploy interactive Streamlit web dashboard with automated verification tests.	Section 12

5. Literature Survey

Author & Year	Methodology	Dataset	Reported Metric	Identified Limitations
Smith et al. (1988)	ADAP Perceptron	Pima Indians (768)	Acc: 76.0%	Early neural model; no systematic missing value imputation; high false negatives.
Kavakiotis et al. (2017)	Systematic Survey	Clinical Databases	Acc: 70% - 82%	Identified pervasive data leakage and unstandardized evaluation protocols.
Sisodia et al. (2018)	DT, SVM, Naive Bayes	Pima Indians (768)	NB: 76.3%	Treated zero values as valid; omitted neural multi-layer architectures.
Zou et al. (2018)	DT, Random Forest	Hospital Exams	RF: 80.84%	Evaluated only raw features without metabolic interaction terms.
Patil et al. (2021)	MLP & Ensembles	Pima Indians (768)	MLP: 79.1%	Omitted L2 regularization grid search; static decision threshold.

6. Dataset Description & Leak-Free Preprocessing

6.1 Pima Indians Diabetes Database: Comprises N = 768 female records with 8 raw attributes and 1 binary target (Class 0: 500 [65.1%], Class 1: 268 [34.9%]). Biologically impossible zeros in continuous variables were identified: Glucose (5 zeros), Blood Pressure (35 zeros), Skin Thickness (227 zeros), Insulin (374 zeros), BMI (11 zeros).

6.2 Preprocessing Workflow: All imputation medians and scaling parameters (mean, standard deviation) were fitted strictly on the training set (N = 537) to prevent data leakage.

Step	Pipeline Stage	Operation / Formula	Leakage Prevention Principle
1	Zero Detection	Mask zeros in Glucose, BP, Skin, Insulin, BMI -> NaN	Independent sample-level masking
2	Median Imputation	$x_{\text{imputed}} = \text{Median}_{\text{train}}$ if NaN else x	Medians computed strictly on D_{train}
3	Feature Engineering	Compute 16 domain features on clean imputed values	Deterministic deterministic transforms
4	Standardization	$z = (x - \mu_{\text{train}}) / \sigma_{\text{train}}$	Fitted StandardScaler applied to Val/Test

7. Domain-Specific Feature Engineering (24 Features)

Feature Category	Engineered Features (16 Total)	Mathematical Formula / Logic
WHO Adiposity Bins	BMI_Underweight, BMI_Normal, BMI_Overweight, BMI_Obese	1-hot indicators for BMI: <18.5, 18.5-24.9, 25-29.9, >=30.0
ADA Glycemic Stages	Glucose_Normal, Glucose_Prediabetes, Glucose_Diabetes	1-hot indicators for Glucose: <100, 100-125, >=126 mg/dL
Age Cohorts	Age_Young, Age_Middle, Age_Senior	1-hot indicators for Age: <30, 30-50, >50 years
Metabolic Interaction	Insulin_Resistance_Proxy	$(\text{Glucose} * \text{Insulin}) / 405.0$
Metabolic Ratio	Insulin_Glucose_Ratio	$\text{Insulin} / (\text{Glucose} + 1e-5)$
Synergistic Risk	BMI_Age_Interaction	$\text{BMI} * \text{Age}$
Gestational Risk	Pregnancy_Age_Risk	$\text{Pregnancies} / (\text{Age} + 1e-5)$
Log Transforms	Log_Insulin, Log_DPF	$\ln(1 + \text{Insulin})$, $\ln(1 + \text{DPF})$ (reduces skewness)

8. Model Architecture & Mathematical Formulation

8.1 Network Topology: The authoritative production model is a Feed-Forward Multilayer Perceptron (Scikit-Learn MLPClassifier) structured as: **Input (24) -> Hidden Layer 1 (64) -> Hidden Layer 2 (32) -> Output (1)**.

8.2 Trainable Parameters:

- Layer 1 (24 -> 64): $24 * 64 + 64 = 1,600$ parameters
- Layer 2 (64 -> 32): $64 * 32 + 32 = 2,080$ parameters
- Output Layer (32 -> 1): $32 * 1 + 1 = 33$ parameters
- **Total Trainable Parameters = 3,713 weights and biases.**

8.3 Objective Function: Binary cross-entropy with L2 weight decay regularization ($\alpha = 0.001$):

$$L(W, b) = -1/N * \sum [y * \ln(\hat{y}) + (1-y) * \ln(1-\hat{y})] + (\alpha / 2N) * \sum ||W_{ij}||^2$$

9. Experimental Results & Benchmark Performance (N = 116)

Performance evaluated on the untouched test partition (N = 116, 76 Negative, 40 Positive cases):

Algorithm	Accuracy	Sensitivity	Specificity	Precision	F1-Score	ROC-AUC	FN	FP
Logistic Regression	79.31%	72.50%	82.89%	0.6905	0.7073	0.8444	11	13
K-Nearest Neighbors	81.90%	77.50%	84.21%	0.7209	0.7470	0.8954	9	12
Support Vector Machine	83.62%	77.50%	86.84%	0.7561	0.7654	0.8808	9	10
Decision Tree	87.93%	82.50%	90.79%	0.8250	0.8250	0.9021	7	7
Random Forest	88.79%	87.50%	89.47%	0.8140	0.8434	0.9431	5	8
Gradient Boosting	89.66%	85.00%	92.11%	0.8500	0.8500	0.9579	6	6
MLP (Baseline)	81.03%	65.00%	89.47%	0.7647	0.7027	0.8625	14	8
MLP (Optimized)	82.76%	75.00%	86.84%	0.7500	0.7500	0.8681	10	10

10. Multi-Model Visual Diagnostics




Figure 1: Multi-Model ROC Benchmark Curves




Figure 2: Confusion Matrices across Classifiers

11. Discussion of Results & Clinical Trade-offs

11.1 Non-Linear Pattern Learning: Hyperparameter tuning improved the MLP's sensitivity from 65.00% to 75.00% while reducing missed cases (False Negatives) from 14 down to 10. The L2 penalty ($\alpha=0.001$) effectively constrained weight norms, avoiding over-fitting on small sample sizes.

11.2 Decision Threshold Calibration: In diabetes screening, False Negatives present substantial clinical hazards. Because the MLP generates smooth continuous probabilities, the decision threshold can be calibrated dynamically: lowering the threshold to $\tau=0.35$ in the Streamlit application reduces False Negatives from 10 to 4, optimizing early intervention.

12. Conclusion & Future Scope

12.1 Conclusion: A robust, leak-free machine learning pipeline was established for diabetes onset prediction. The optimized Scikit-Learn Multilayer Perceptron demonstrated solid performance (82.76% Accuracy, 0.8681 ROC-AUC) on untouched test data. An enterprise Streamlit Cloud web dashboard was successfully created and deployed.

12.2 Future Scope:

- Validate on multi-center international clinical cohorts (e.g., NHANES, UK Biobank).
- Benchmark against tabular attention models (TabNet, FT-Transformer).
- Integrate explainable AI attribution (SHAP values, Integrated Gradients).
- Export quantized ONNX models for low-power edge mobile screening devices.

13. Key References

- Smith, J. W., et al. (1988). Using the ADAP learning algorithm to forecast the onset of diabetes mellitus. *Proc. Symp. Comput. Appl. Med. Care*, 261–265.
- American Diabetes Association. (2024). Standards of Medical Care in Diabetes—2024. *Diabetes Care*, 47(Suppl. 1), S1–S343.
- Kavakiotis, I., et al. (2017). Machine learning and data mining methods in diabetes research. *Comput. Struct. Biotechnol. J.*, 15, 104–116.
- Sisodia, D., & Sisodia, D. S. (2018). Prediction of diabetes using classification algorithms. *Procedia Computer Science*, 132, 1578–1585.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.